

The “Shiva runtime environment”



Advancing the programmability of the Linux
process address space

Ryan “elfmaster” O’Neill: ryan@bitlackeys.org

What is Shiva exactly?

- A weird machine: programmable runtime engine
- Custom dynamic linker:
 - Loads ELF micro-programs into the address space
- Shiva provides a rich API for program transformation, tracing, hooking, and debugging.
- Think “LKM’s (Loadable kernel modules) for userland processes”.
- A programmable debugging engine that does ***NOT require the ptrace system call***



What can Shiva be used for?

- Building debuggers and tracers
- Designing process Security modules
 - Anti-exploitation, Sandboxing, Process hardening
- In-memory fuzzing harnesses
- Software profiling
- Virus detection engines
- Malware unpackers
- Hot-patching code and data
- Userland-Rootkit detection & disinfection
- Extremely fast instrumentation, hooking, and process injection of all types (*Without the use of SYS_ptrace*).



From a blackhat perspective

- Shiva can be used to crack software
- Design powerful in-process rootkits
- Perform Modular virus infection: “Preloading the linker for fun and profit”: <https://tmpout.sh/2/6.html>
- Create ELF binary packers

How does Shiva run?

- Shiva loads and links “**Shiva modules**” at runtime
- Shiva loads the module “/opt/shiva/modules” directory
 1. Shiva can be invoked directly as an executable
 - *Shiva works as a userland-exec to load and map the target executable into its own process address space (See grugqs work on ul_exec: https://grugq.github.io/docs/ul_exec.txt)*
 - *i.e. `./shiva /bin/test_prog`*
 2. Shiva can be invoked indirectly as a runtime interpreter (*i.e. A program declares shiva as the dynamic linker*)
 - *The kernel loads the target executable and loads Shiva as the interpreter*
 - *Good for writing modules that are linked at every single execution*
 - *i.e. `/bin/test_prog`*



Shiva ran directly on program ./test

- Shiva is loading a module: `pltgot_hook.o`
- This module hooks the PLT function '`puts`' and modifies the string output

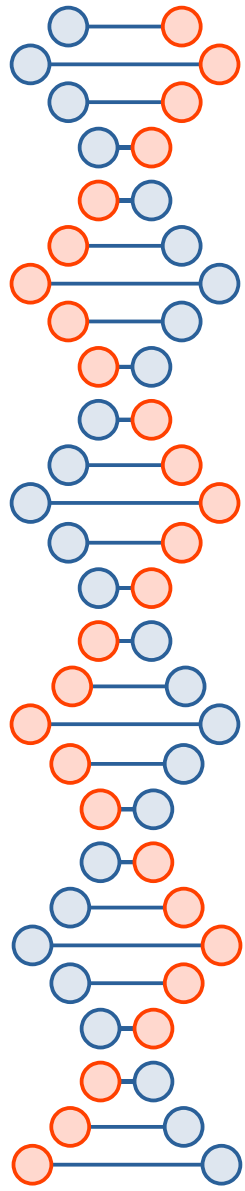
```
elfmaster@esoteric-labs:~/git/shiva$ ./test abc
Hello World
Hello World
Hello World
Hello World
Hello World
abc
elfmaster@esoteric-labs:~/git/shiva$ ./shiva test abc
hijacked your string: 'Hello World'
hijacked your string: 'Hello World'
hijacked your string: 'Hello World'
hijacked your string: 'Hello World'
hijacked your string: 'Hello World'
hijacked your string: 'abc'
elfmaster@esoteric-labs:~/git/shiva$
```

Shiva interpreter path

- The INTERP segment contains the path to 'shiva' instead of "ld-linux.so".

```
Elf file type is DYN (Position-Independent Executable file)
Entry point 0x1080
There are 14 program headers, starting at offset 64

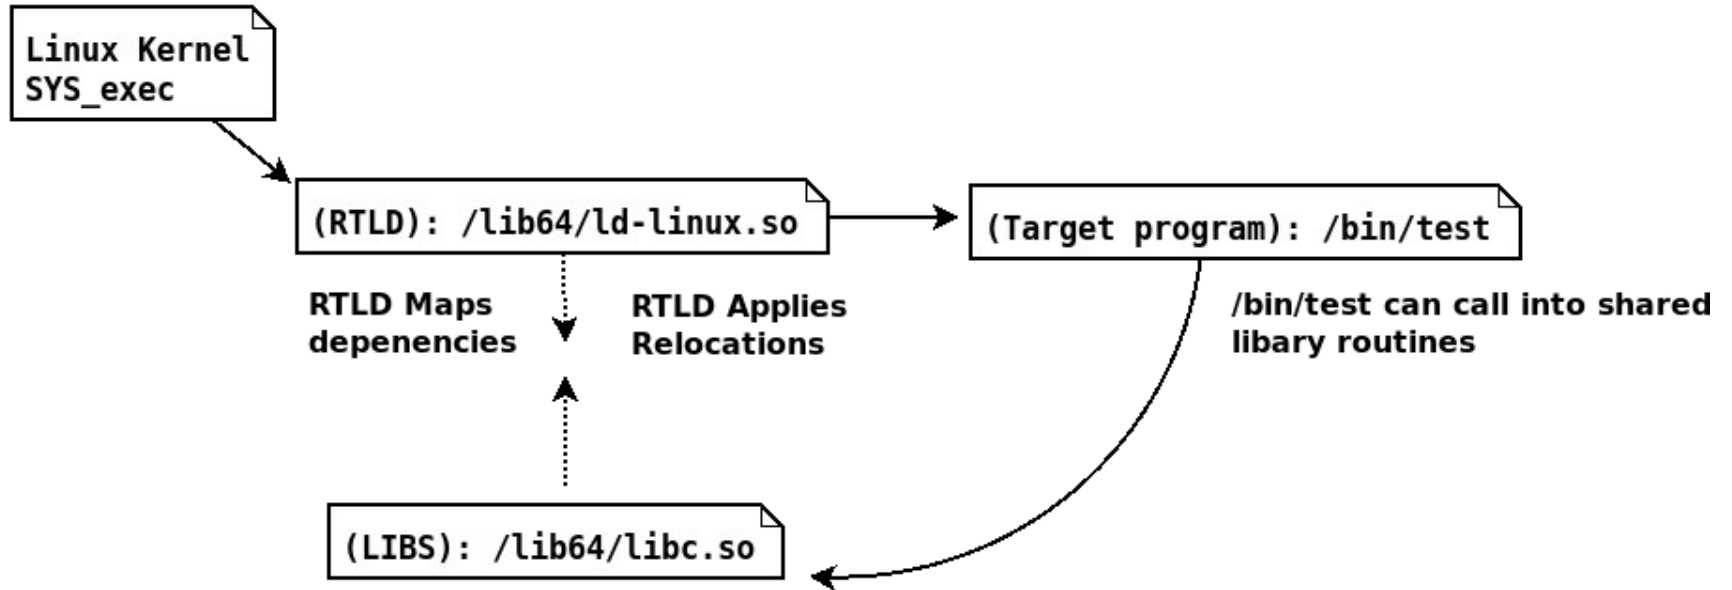
Program Headers:
Type           Offset             VirtAddr           PhysAddr
               FileSiz              MemSiz              Flags  Align
PHDR           0x0000000000000040 0x0000000000000040 0x0000000000000040
               0x0000000000000310 0x0000000000000310  R      0x8
INTERP         0x0000000000000350 0x0000000000000350 0x0000000000000350
               0x0000000000000020 0x0000000000000020  R      0x1
                [Requesting program interpreter: /home/elfmaster/git/shiva/shiva]
LOAD           0x0000000000000000 0x0000000000000000 0x0000000000000000
```



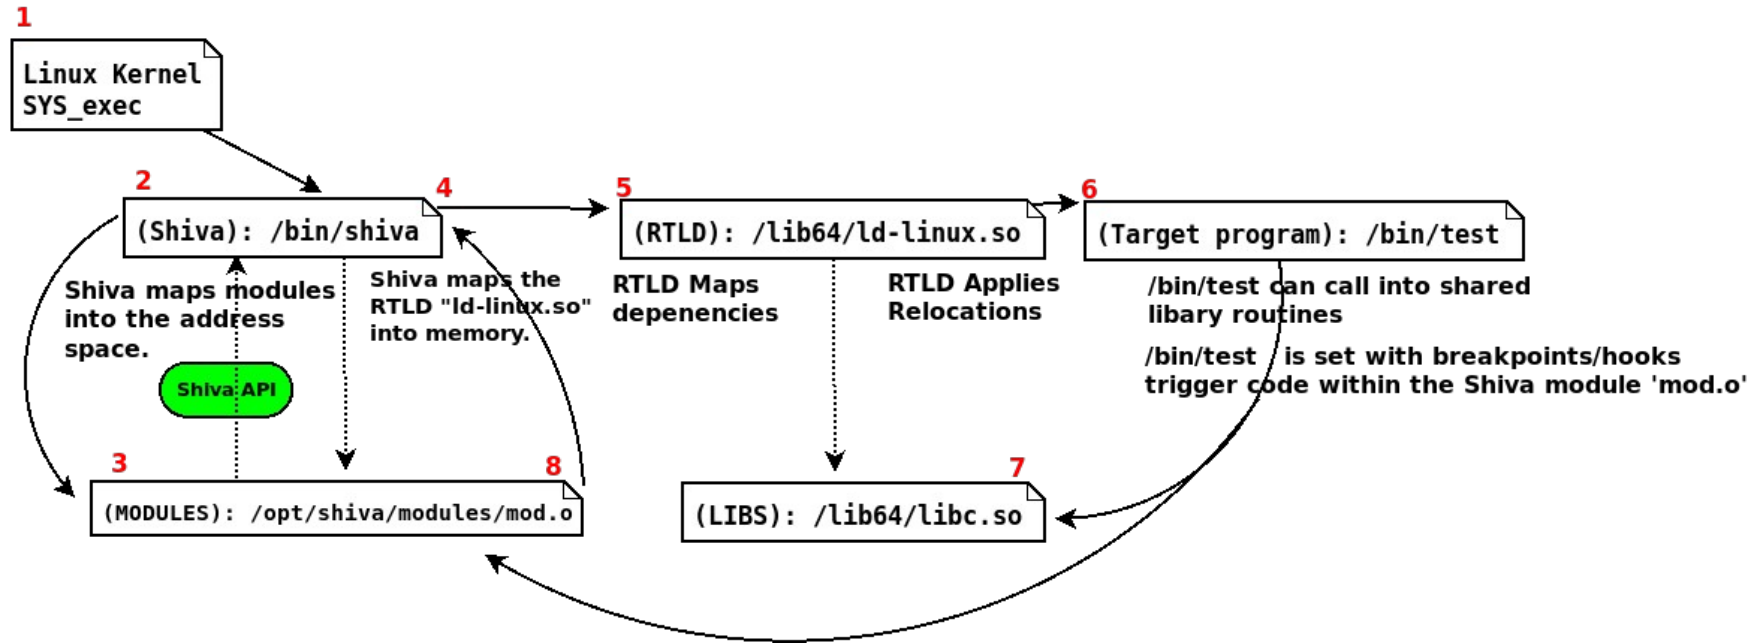
Shiva running as the interpreter for ./test

```
elfmaster@esoteric-labs:~/git/shiva$ ./test abc  
hijacked your string: 'Hello World'  
hijacked your string: 'Hello World'  
hijacked your string: 'Hello World'  
hijacked your string: 'Hello World'  
hijacked your string: 'Hello World'  
hijacked your string: 'abc'  
elfmaster@esoteric-labs:~/git/shiva$
```


Process execution flow of a normal dynamically linked ELF



Process execution flow with Shiva as an interpreter.





Understanding Shiva modules

- Shiva modules are loaded from:
`"/opt/shiva/modules"`
- Similar to LKM (Loadable kernel module):
 - Shiva modules are ELF relocatable object
 - Shiva modules must have an initialization routine
`int shakti_main(shiva_ctx_t *ctx)`

Shiva Module: context struct

- The context struct “**shiva_ctx_t**”
- Passed into the modules init function
- Contains:
 - Register state
 - Process-state
 - Thread-state
 - Control flow data
 - Signal data
 - Shiva-Trace API specific data structures
 - Full access to **elfobj_t** of target executable

Shiva modules: Linking

- Shiva modules use the Shiva_trace API for program instrumentation
- Modules are written in C and have access to:
 - `musl-libc`
 - `libelfmaster`
 - `Libudis86`
 - `Shiva-Trace API`
- Shiva loader handles loading and linking of modules
- The module is an **ELF ET_REL** (relocatable object file)
- Sections are mapped appropriately into memory segments: Text, and Data



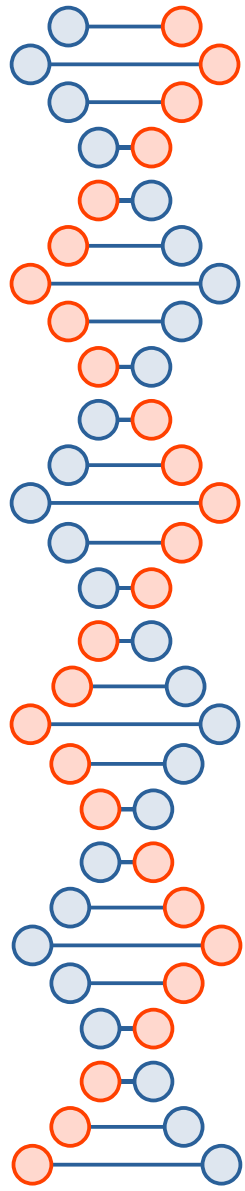
Shiva modules execute in-process

- Shiva modules are executed **in-process**.
- **In-process** means:
 - The Shiva runtime engine is in the same process as the target program
 - The module also executes within the same process address space
- Instrumentation and hooking is extremely fast
- Module has an ***initialization function***
- Module init function sets hooks and breakpoints within the target program
- Handler functions are callbacks for hooks and breakpoints



Understanding Shiva modules: **gcc code models**

- Shiva modules are currently compiled with GCC and must be compiled with a large code model (i.e.
`gcc -c -mcmodel=large`)
- A large code model is necessary to handle module relocation offsets that exceed 4GB
- On 64bit system: The Shiva executable, the Shiva module, and the target executable will be mapped to addresses that are vastly far apart
- In some instances we can and must use a small code model for modules (More on this later)



A security module. Prevent command injection via system function

```
#include "../shiva.h"

int n_system(const char *s)
{
    printf("s: %s\n", s);
    if (strstr(s, ";") != NULL || strstr(s, "|") != NULL) {
        printf("Detected possible OS command injection attack '%s'\n", s);
        abort();
    }
    return system(s);
}

int
shakti_main(shiva_ctx_t *ctx)
{
    bool res;
    shiva_error_t error;

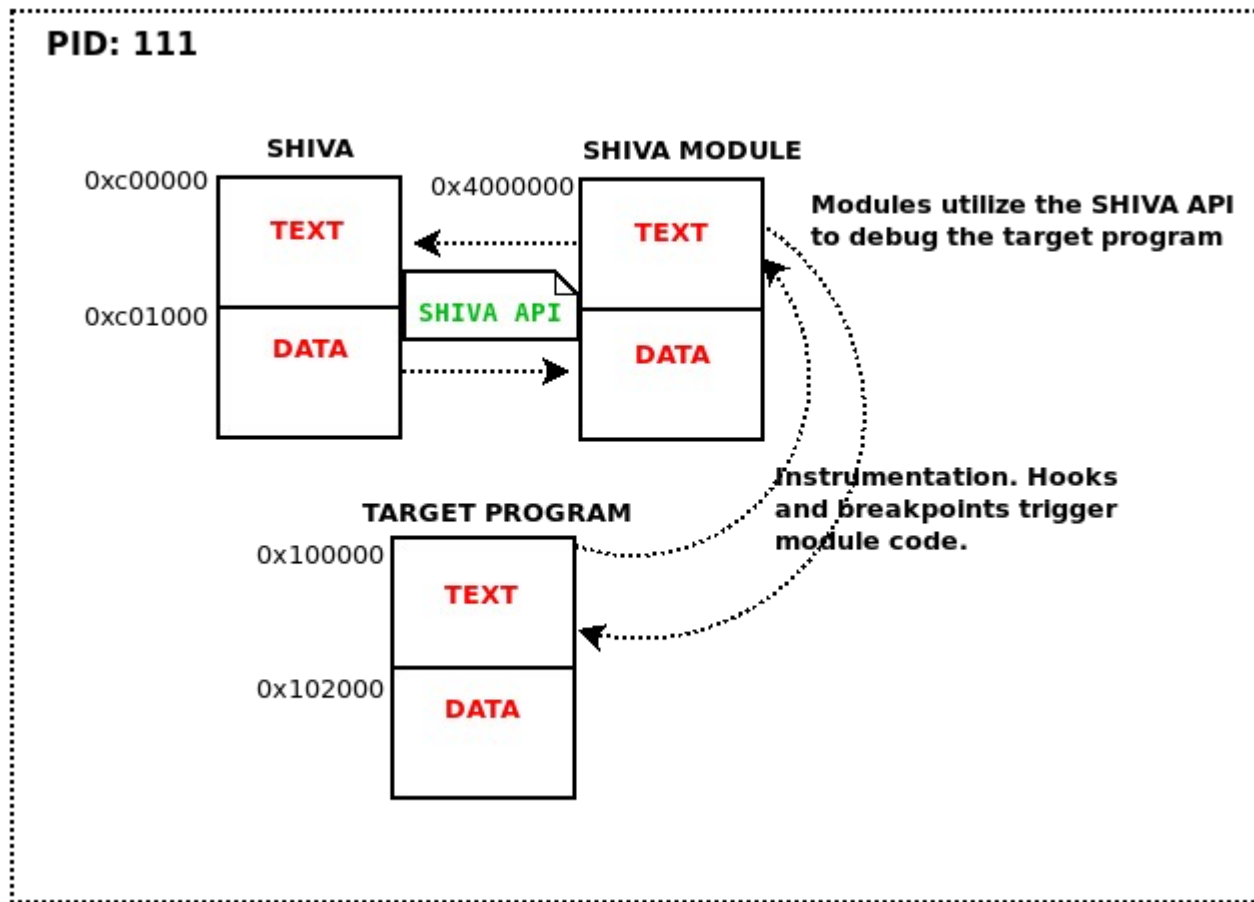
    res = shiva_trace(ctx, 0, SHIVA_TRACE_OP_ATTACH,
        NULL, NULL, 0, &error);
    if (res == false) {
        printf("shiva_trace failed: %s\n", shiva_error_msg(&error));
        return -1;
    }
    res = shiva_trace_register_handler(ctx, (void *)&n_system,
        SHIVA_TRACE_BP_PLTGOT, &error);
    if (res == false) {
        printf("shiva_register_handler failed: %s\n",
            shiva_error_msg(&error));
        return -1;
    }
    res = shiva_trace_set_breakpoint(ctx, (void *)&n_system,
        0, "system", &error);
    if (res == false) {
        printf("shiva_trace_set_breakpoint failed: %s\n", shiva_error_msg(&error));
        return -1;
    }
    return 0;
}
```



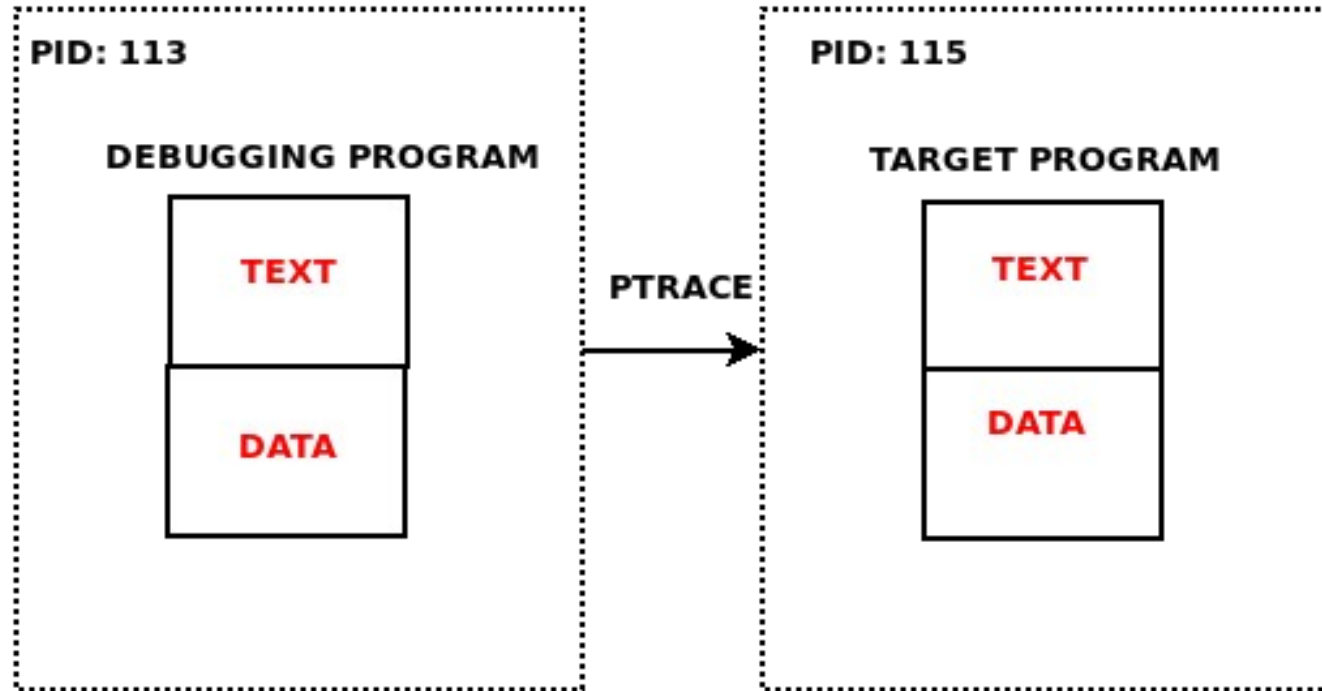

The Shiva API (shiva_trace)

- Shiva modules are written in C
- The Shiva_trace API can: **access, read, transform and instrument the process image.**
- Rich API for process transformation
- Excellent for writing debuggers, tracers, security modules, etc.
- In very early development phase

Shiva debugging model



PTRACE Debugging model





Shiva-Trace vs. PTRACE

- The Shiva-Trace API does not rely on the **PTRACE** system call
- Shiva-Trace is effective against hardened binaries and anti-debugging
- Shiva-Trace is implemented as an “in-process” programmable runtime engine.
- **PTRACE** on the other hand, executes from a separate process context.
- **PTRACE**: Slow. And is useless against many Linux anti-debugging techniques.
- **PTRACE** operations send a **SIGSTOP** to the target process.
- **PTRACE** is a system call and requires a context switch, whereas Shiva_trace is implemented purely in userland and less expensive.



Shiva Trace p3.

- shiva_trace can cleverly instrument the target program at runtime
- shiva_trace uses the terminology “**Hooks**” and “**Breakpoints**” interchangeably.
- Many breakpoints are really just various types of hooks
- shiva_trace can defeat anti-debugging techniques
- Powerful for debugging hostile binaries

Shiva API basics: shiva_trace()

```
/*
 * shiva_trace lets the module attach to the process, read and write to the process,
 * get and set the register state.
 */
bool shiva_trace(shiva_ctx_t *, pid_t, shiva_trace_op_t, void *, void *, size_t, shiva_error_t *);

typedef enum shiva_trace_op {
    SHIVA_TRACE_OP_CONT = 0,
    SHIVA_TRACE_OP_ATTACH,
    SHIVA_TRACE_OP_POKE,
    SHIVA_TRACE_OP_PEEK,
    SHIVA_TRACE_OP_GETREGS,
    SHIVA_TRACE_OP_SETREGS,
    SHIVA_TRACE_OP_SETFPREGS,
    SHIVA_TRACE_OP_GETSIGINFO,
    SHIVA_TRACE_OP_SETSIGINFO
} shiva_trace_op_t;
```



Shiva-Trace API - Breakpoints

- Breakpoints are registered to a handler function
- Similar to Linux Kprobes, but for userland
- Multiple breakpoints can be set for a single handler
- Some breakpoints are hooks
- Some breakpoints are signal driven
- A complex series of in-process trampolines



Shiva breakpoints and signals

- We do not use PTRACE to catch signals
- Our debugging module exists within the same PID as the program it's debugging
- Signal driven breakpoints (i.e. `int3`) register a handler internally with `sigaction()`
- Shiva internally hooks `signal()` and `sigaction()` in the target program to prevent the target from overwriting breakpoint handlers



A look at all Shiva breakpoint types

```
typedef enum shiva_trace_bp_type {  
    SHIVA_TRACE_BP_JMP = 0,  
    SHIVA_TRACE_BP_CALL,  
    SHIVA_TRACE_BP_INT3,  
    SHIVA_TRACE_BP_SEGV,  
    SHIVA_TRACE_BP_SIGILL,  
    SHIVA_TRACE_BP_TRAMPOLINE,  
    SHIVA_TRACE_BP_PLTGOT  
} shiva_trace_bp_type_t;
```

Setting hooks and breakpoints

- Breakpoints/hooks are registered to a handler function

```
res = shiva_trace_register_handler(ctx, (void *)&bp_handler,  
    SHIVA_TRACE_BP_INT3, &error);  
if (res == false) {  
    printf("shiva_register_handler failed: %s\n",  
        shiva_error_msg(&error));  
    return -1;  
}  
if (elf_symbol_by_name(&ctx->elfobj, "main", &symbol) == false) {  
    printf("couldn't find symbol main in target\n");  
    return -1;  
}  
res = shiva_trace_set_breakpoint(ctx, (void *)&bp_handler,  
    symbol.value + ctx->ulexec.base_vaddr, NULL, &error);  
if (res == false) {  
    printf("shiva_trace_set_breakpoint failed: %s\n", shiva_error_msg(&error));  
    return -1;  
}
```

Basic handler function

```
void  
bp_handler(int sig, siginfo_t *si, void *data)  
{  
  
    ucontext_t *uctx = (ucontext_t *)data;  
  
    printf("Breakpoint at %#lx\n", uctx->uc_mcontext.gregs[REG_RIP] - 1);  
    return;  
}
```



SHIVA_TRACE_BP_JMP

- **Breakpoint type:** Control flow hook
- Set on 5 byte jmp instructions
- The jmp offset is replaced by an offset to your registered handler function
- When done, the handler jumps back to the original target.
- **SHIVA_TRACE_LONGJMP_RETURN** macro: rewinds the stack and jmp's to the original offset
- Defeats breakpoint detection with anti-debugging

SHIVA_TRACE_BP_CALL

- **Breakpoint type:** Control flow hook
- Set on 5 byte immediate call instructions
- Original call target is replaced with offset to registered handler
- The registered handler returns after calling the original function (If desired)
- Extremely fast way to build a function tracer: (See `modules/examples/func_tracer.c`)
- Only works when executing Shiva directly (Not as an interpreter) due to required small-code model.
- Defeats breakpoint detection



SHIVA_TRACE_BP_TRAMPOLINE

- **Breakpoint type:** Control flow hook
- This type of breakpoint inserts a trampoline on the entry point of the function
 - `movq $target_addr, %rax`
`jmp *%rax`
- Excellent for hooking any function locally or globally within the address space
- Trampoline function restores the original code bytes before calling the original function
- Not currently thread-safe. Memcpy's that place and restore hooks are ***not atomic operations***
- Works well in a large code model whereas ***call hooks*** do not.



SHIVA_TRACE_BP_PLTGOT

- Hook any shared library function within the PLT/GOT
- Modifies the `.got.plt[entry]` of the given symbol
- Was tricky to implement. The RTLD wants to patch over our GOT patch. (i.e. with **DT_BINDNOW**)
- Very low overhead, and thread safe
- Handler is usually a hook-replacement function for the function being hooked
- A single handler can also be used for multiple PLTGOT hooks. (*i.e. Any PLT call triggers the same handler*)



SHIVA_TRACE_BP_INT3

- **Breakpoint type:** Signal driven
- Traditional `int3` breakpoint uses `0xcc` opcode
- Causes kernel to deliver **SIGTRAP**
- The breakpoint handler is essentially a signal handler
- Uses `struct ucontext` to get register state in this case, as any signal handler would
- If the target program already had a **SIGTRAP** handler: ours replaces it, and optionally can invoke the original handler



SHIVA_TRACE_BP_SIGILL

- **Breakpoint type:** Signal driven
- This breakpoint inserts a two byte **ud2** (undefined instruction) to trigger **SIGILL**. `"\x0f\x0b\"`
- Breakpoint is registered to your signal handler function
- Can be helpful against hardened binaries that perform anti-debugging on standard breakpoints



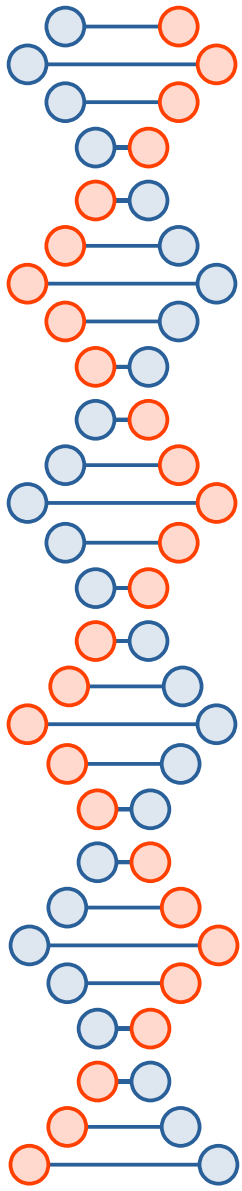
SHIVA_TRACE_BP_SEGV

- **Breakpoint type:** Signal driven
- Breakpoint code:
 - `push <invalid_offset>(%rip)`
- Shiva_trace API inserts an instruction with an invalid offset or address
- Causes an invalid memory dereference, and thus a **SIGSEGV**
- The handler function is a signal handler that catches the **SEGV**, and executes your custom code



Reading and writing to process memory

- The `shiva_trace()` API call
- **SHIVA_TRACE_OP_PEEK**
 - Allows the module to read anywhere in process memory
- **SHIVA_TRACE_OP_POKE**
 - Allows the module to write anywhere in process memory (Except to the Shiva code itself)



Getting and setting register state

- The shiva_trace API call
- **SHIVA_TRACE_OP_GETREGS**
 - Get the current register values
- **SHIVA_TRACE_OP_SETREGS**
 - Set the register values



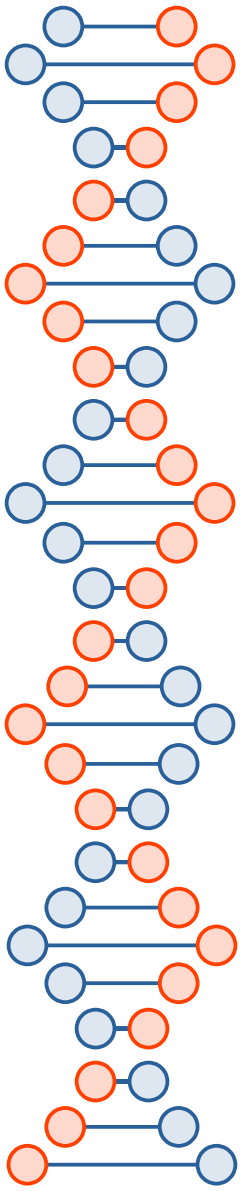
In-process debugging tools with Shiva

- When running Shiva directly it can be used for creating powerful debugging tools
- “In-process” debugging technology is extremely fast
- Less context switching, more in-process trampolines
- **modules/examples/func_tracer.c** is a module for function tracing
- Sets a **SHIVA_TRACE_BP_CALL** hook on every immediate call
- Registered handler function simply prints each function name being called



Shiva module: func_tracer.o

- func_tracer.c
- Only 84 lines of code
- Much faster than tracers that are written using **PTRACE**
- Benchmarks demonstrating performance:
- Command **./test**:
 - **Real 0m0.052s**
- Tracing **./test** with Shiva:
 - **Real 0m0.155s**
- Tracing **./test** with **/bin/ltrace**
 - **Real 0m3.582**



Control flow integrity: ret2PLT protection

- **modules/examples/plt_cfi.c**
- Only 109 lines of code
- Forward edge / Backward edge CFI
- Sets a **SHIVA_TRACE_BP_PLTGOT** hook on every PLT call
- **shiva_trace_set_breakpoint()** installs hooks/breakpoints and stores CFI data in breakpoint struct
- Every PLT call is redirected to our modules handler:
int plt_handler(void *arg)



Ret2plt handler function

- **`int plt_handler(void *arg)`**
- Checks the return address
- Prevents arbitrary calls to the PLT
- A ret2PLT attack will fail against this



Mitigate OS command injection

- `modules/examples/os_inject.o`
- Only 48 lines of code
- `SHIVA_TRACE_BP_PLTGOT` hook on `system()` function
- Replaced with secure function `n_system()`
- `n_system()` sanitizes the command string



Crack software

- `./crackme`
- Shiva module to crack serial number validator
- **`modules/examples/crackme_bypass.o`**
- Redirects: **`int check_serial(char *serial)`**
- To a modified version of the function
- Uses a **`SHIVA_TRACE_BP_TRAMPOLINE`** breakpoint



Shiva for debugging hardened binaries

- There is a long history of anti-debugging techniques
- PTRACE system call is easily foiled
- PTRACE fails on PaX hardened binaries
 - `mprotect` restrictions prevent `ptrace(PTRACE_POKETEXT, ...)` ;
- Shiva maps the hostile/hardened program into it's own address space
- In-process debugging can be very effective for Reverse engineering tough problems



Tracing a hardened binary

- **modules/examples/func_tracer.o**
- Capable of tracing binaries with anti-debugging
- **test_antidebug.c**: Simple anti-debugging techniques
 - **ptrace(PTRACE_TRACEME, 0, 0, 0, 0);**
 - **prctl(PR_SET_DUMPABLE, 0, 0, 0, 0);**



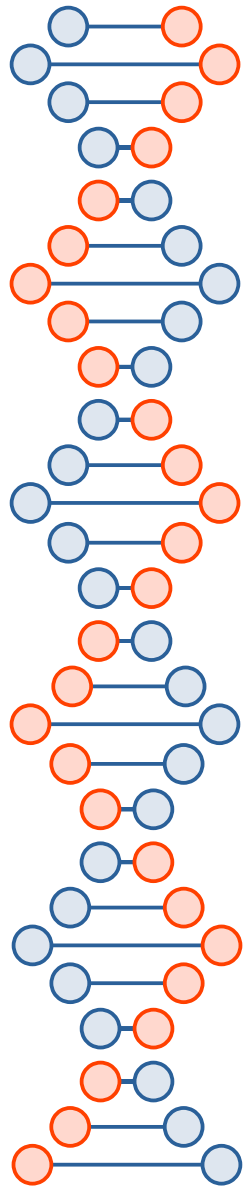
Shiva works on stripped ELF binaries

- Corrupted meta-data
- Missing symbol tables
- Missing section header table
- Shiva uses libelfmaster:
<https://github.com/elfmaster/libelfmaster>
- Forensic reconstruction of stripped binaries:
 - Symbol table reconstruction
 - Section header table reconstruction



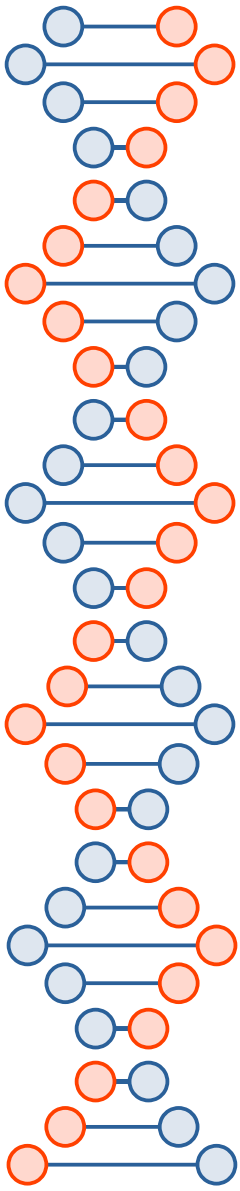
Tracing a stripped binary

- ltrace fails to get any tracing info:
 - Requires section headers and symbol tables
- Our Shiva module for function tracing succeeds
 - Forensically reconstructs section and symbol data



Current limitations

- Support for multi-threaded applications
- Shiva doesn't work on ELF's with **-fcf-protection**
- Shiva can load only one module at a time
- Shiva only works with **x86_64 ET_DYN** binaries
 - Will add support for **x86_64 ET_EXEC**
- Running Shiva as an interpreter requires a large code model for all modules
 - Some Shiva features require a small code model (i.e. **SHIVA_TRACE_BP_CALL** breakpoints)



Future plans

- Open source the technology
- Continue to develop and expand the usability and functionality
- **shiva-gcc** wrapper for modules
- Design a suite of security modules for process hardening
- Handle loading multiple modules
- Handle threads

A note on handling threads

- How do we handle threads without using SYS_ptrace?
- TLS resolver hooks
- pthread_create hooks
- SYS_clone hooks
- Shared memory segments



Related works

- Linux kprobe instrumentation
 - <http://phrack.org/issues/67/6.html>
- Preloading the linker for fun and profit
 - <https://tmpout.sh/2/6.html>
- Embedded ELF debugging (by Mayhem & ERESI)
 - <http://phrack.org/issues/63/9.html>
- Userland exec implementation (By the Grugq:)
 - https://grugq.github.io/docs/ul_exec.txt



Questions?

- Ryan@bitlackeys.org
- [Bitlackeys.org](https://bitlackeys.org)
- <https://github.com/elfmaster/shiva>